

Logistic Regression

- Logistic Regression is a statistical model used for binary classification tasks. It predicts the probability that an input belongs to one of two possible classes, typically denoted as 0 and 1.
- Logistic Regression uses the sigmoid (logistic) function to map the linear combination of input features to a probability value between 0 and 1. $P(Y=1|X) = 1 / (1 + e^{-(b_0 + b_1x_1 + b_2x_2 + \dots + b_nx_n)})$

$$= 1 / (1 + e^{-z})$$
 - The model minimizes a loss function (e.g., log loss) to find the best fit.

```
In [1]: import seaborn as sns
import pandas as pd
import numpy as np
```

```
In [2]: df = sns.load_dataset('iris')
df.head()
```

```
Out[2]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
In [3]: df['species'].unique()
```

```
Out[3]: array(['setosa', 'versicolor', 'virginica'], dtype=object)
```

```
In [4]: df.isnull().sum()
```

```
Out[4]: sepal_length
sepal_width
petal_length
petal_width
species
```

```
In [5]: # Only taking two classes
df = df[df['species']=='setosa']
```

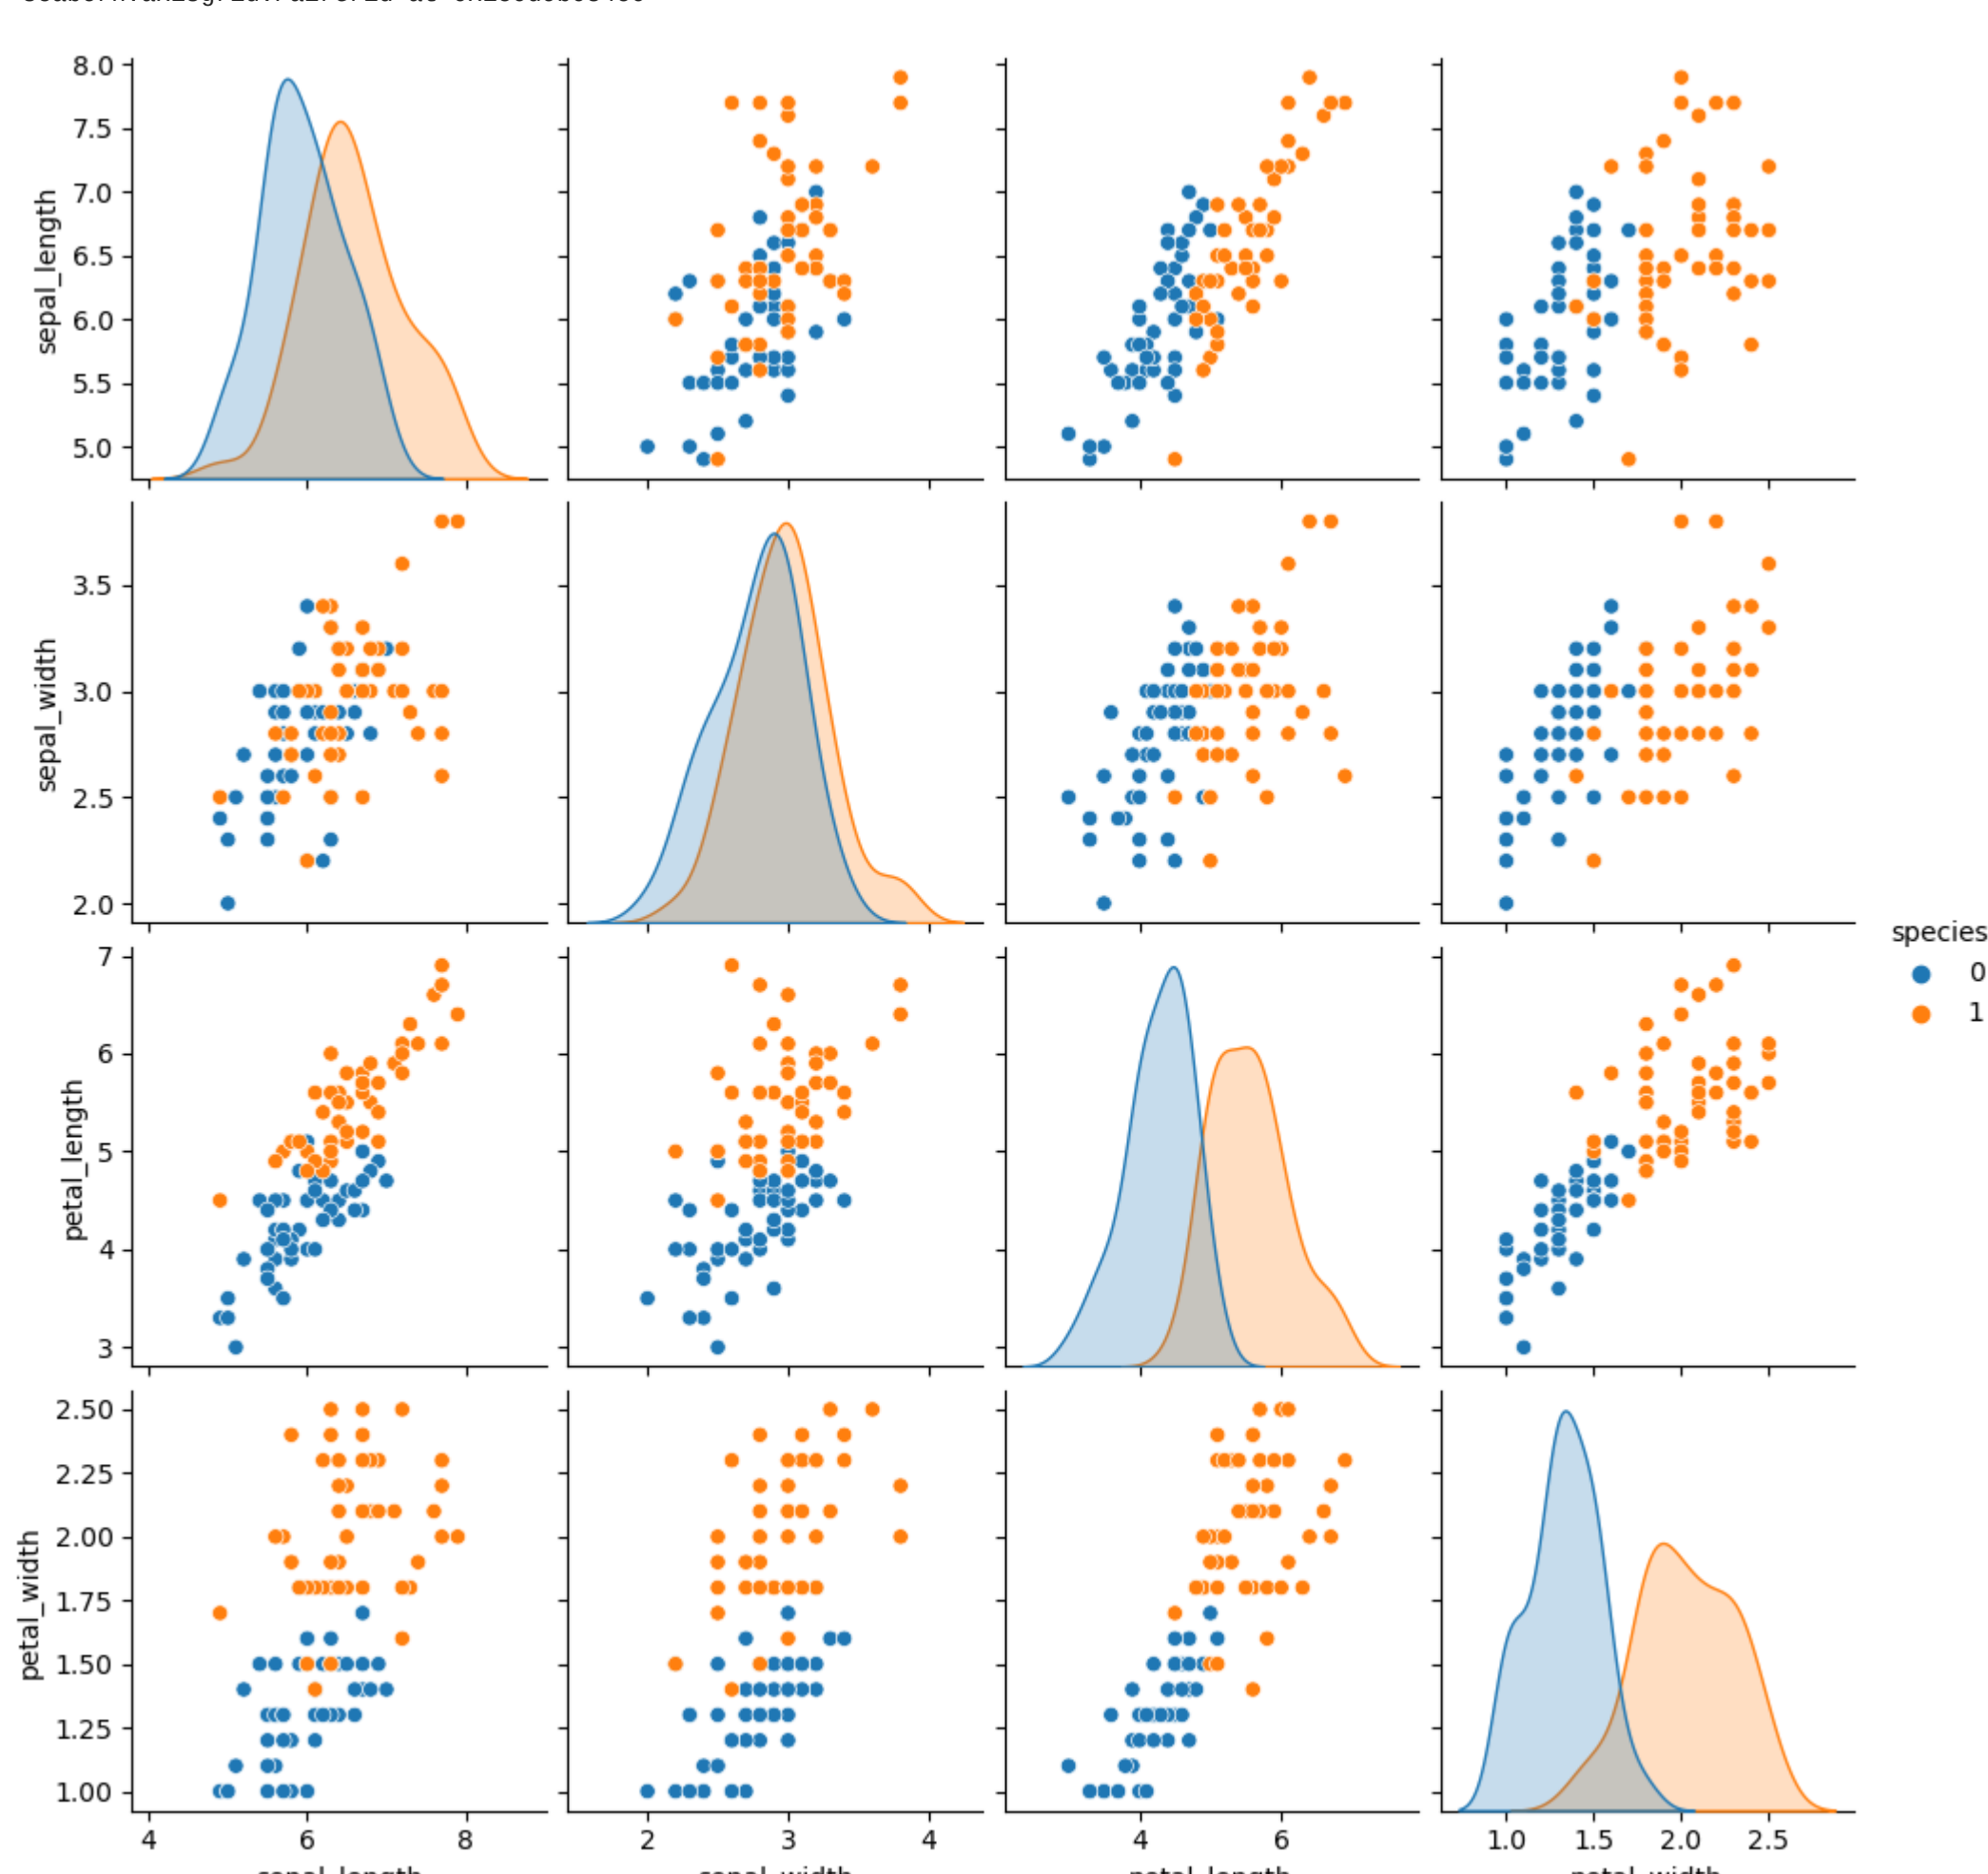
```
Out[6]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
50	7.0	3.2	4.7	1.4	versicolor
51	6.4	3.2	4.5	1.5	versicolor
52	6.9	3.1	4.9	1.5	versicolor
53	5.5	2.3	4.0	1.3	versicolor
54	6.5	2.8	4.6	1.5	versicolor

```
Out[8]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
50	7.0	3.2	4.7	1.4	0
51	6.4	3.2	4.5	1.5	0
52	6.9	3.1	4.9	1.5	0
53	5.5	2.3	4.0	1.3	0

```
sns.pairplot(df, hue = 'species')
```



```
Out[10]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
sepal_length	1.000000	0.553855	0.828479	0.593709	0.494305
sepal_width	0.553855	1.000000	0.519802	0.566203	0.308080
petal_length	0.828479	0.519802	1.000000	0.823348	0.786424
petal_width	0.593709	0.566203	0.823348	1.000000	0.828129

```
In [11]: # Split dataset into independent and dependent features
X = df.iloc[:, :-1]
y = df.iloc[:, -1]
```

	sepal_length	sepal_width	petal_length	petal_width
50	7.0	3.2	4.7	1.4
51	6.4	3.2	4.5	1.5
52	6.9	3.1	4.9	1.5
53	5.5	2.3	4.0	1.3
54	6.5	2.8	4.6	1.5
...
145	6.7	3.0	5.2	2.3
146	6.3	2.5	5.0	1.9
147	6.5	3.0	5.2	2.0
148	6.2	3.4	5.4	2.3
149	5.9	3.0	5.1	1.8

```
In [13]: y
Out[13]: 50      0
          51      0
          52      0
          53      0
          54      0
          ..
          145     1
          146     1
          147     1
          148     1
          149     1
```

```
In [14]: from sklearn.model_selection import train_test_split
```

```
In [15]: from sklearn.linear_model import LogisticRegression
```

```
Classifier = LogisticRegression()

In [16]: from sklearn.model_selection import GridSearchCV
```

```
parameter = {'penalty': ['l1', 'l2', 'elasticnet'], 'C': [1, 2, 3, 4, 5, 6, 10, 20, 30, 40, 50], 'max_iter': [100, 200, 300, 400, 500, 600, 700, 800, 900, 1000]}
```

```
[In [17]: classifier_regressor = GridSearchCV(classi
```

[illegible]

```
Out[18]: warnings.warn(
  ▶ GridSearchCV
  ▶ estimator: LogisticRegression
    ▶ LogisticRegression
```

```
In [19]: print(classifier_regressor.best_params_)
```

```
{'C': 1, 'max_iter': 100, 'penalty': 'l2'}
```

0.9733333333333334

```
In [21]: #Prediction
y_pred = classifier_regressor.predict(X_test)
```

```
# Accuracy_Score
from sklearn.metrics import accuracy_score, classification_report

score = accuracy_score(y_pred, y_test)
print(score)

0.92

print(classification_report(y_pred, y_test))
```

	precision	recall	f1-score	support
0	0.93	0.93	0.93	14
1	0.91	0.91	0.91	11
accuracy	0.92	0.92	0.92	25